

# Operating System (CSC 3150)

## Tutorial 3

*Lele Li:*  
*221019053@link.cuhk.edu.cn*

# Target

In this tutorial, we will practice **Pthread** programming using c/c++.

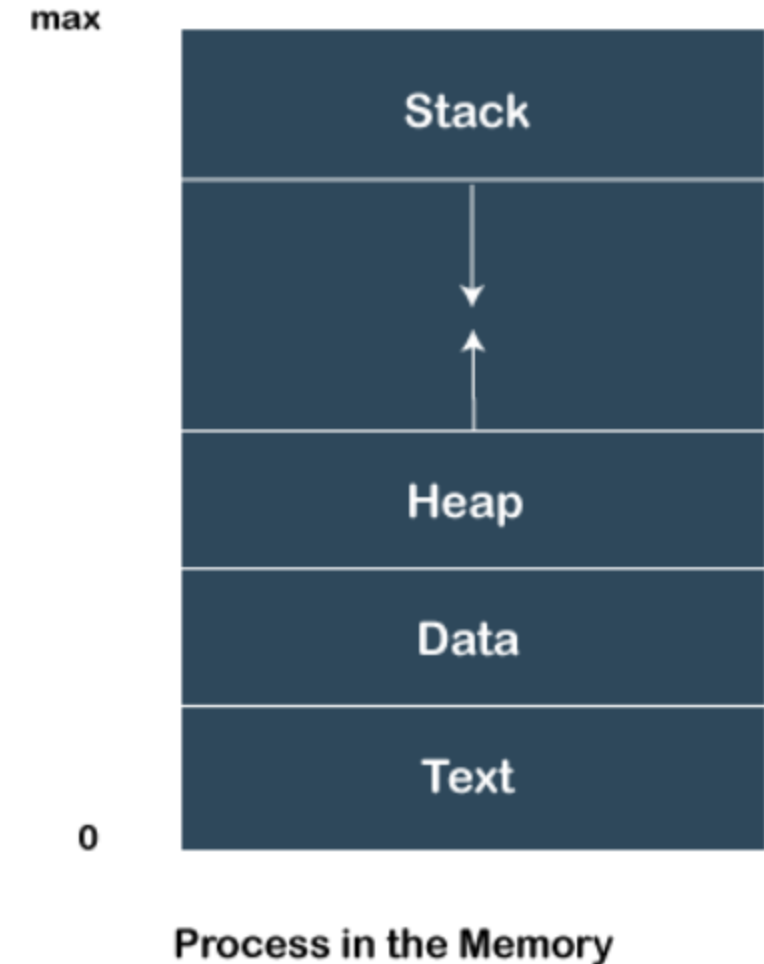
- Process
- Thread
- Pthread creation
- Pthread termination
- Pthread join
- Pthread mutex
- Pthread condition

## What is Process?

- A process is **an instance of a program** that is being executed.

### Features of Process

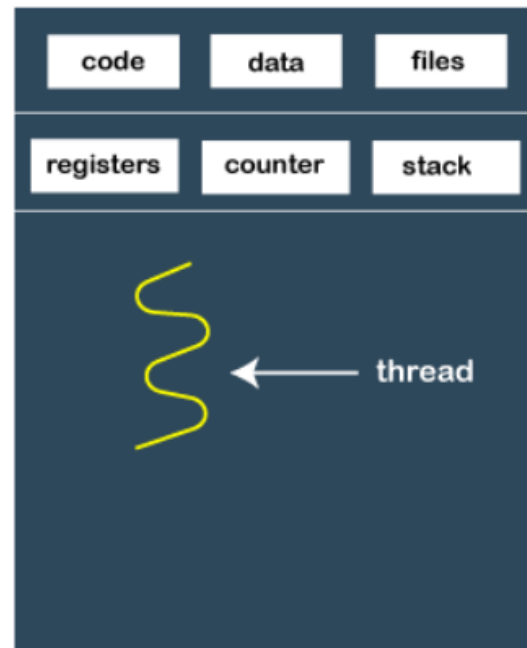
- Each time we create a process, we need to make a separate system call for each process to the OS. The **fork()** function creates the process.
- **Each process exists within its own address or memory space.**
- Each process is **independent and treated as an isolated process by the OS.**
- Processes need IPC (Inter-process Communication) in order to communicate with each other.



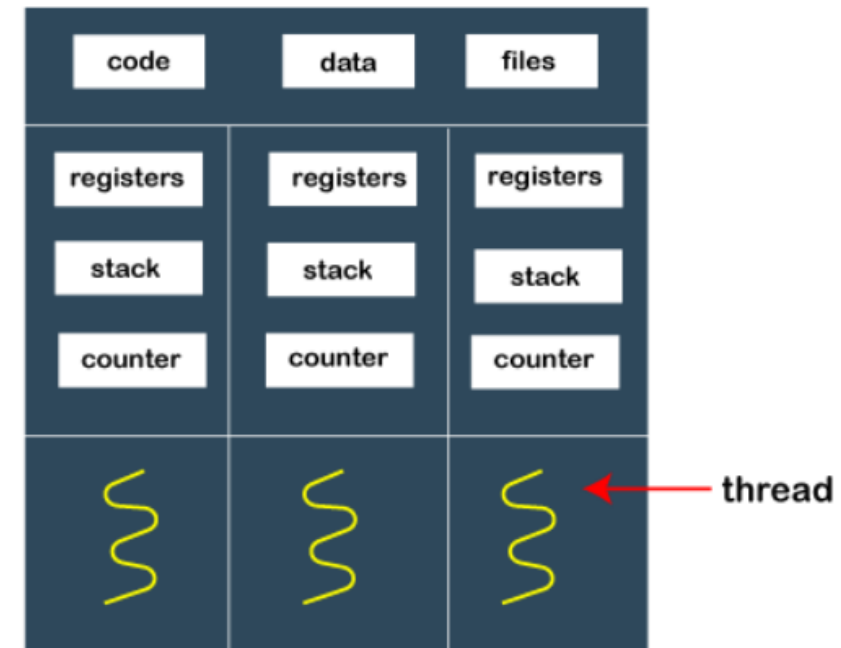
## What is a thread?

A thread is a single sequence stream within a process. Because threads have some of the properties of processes, they are sometimes called *lightweight processes*.

- Threads are not independent from each other unlike processes. As a result, **threads shares with other threads their code section, data section and OS resources like open files and signals**. But, like processes, a thread has its own **program counter (PC), a register set, and a stack space**.
- Threads use and exist within these process resources. They are able to be scheduled by the operating system and run as independent entities within a process.
- A process can have multiple threads**, all of which share the resources within a process and all of which execute within the same address space.



Single-threaded process



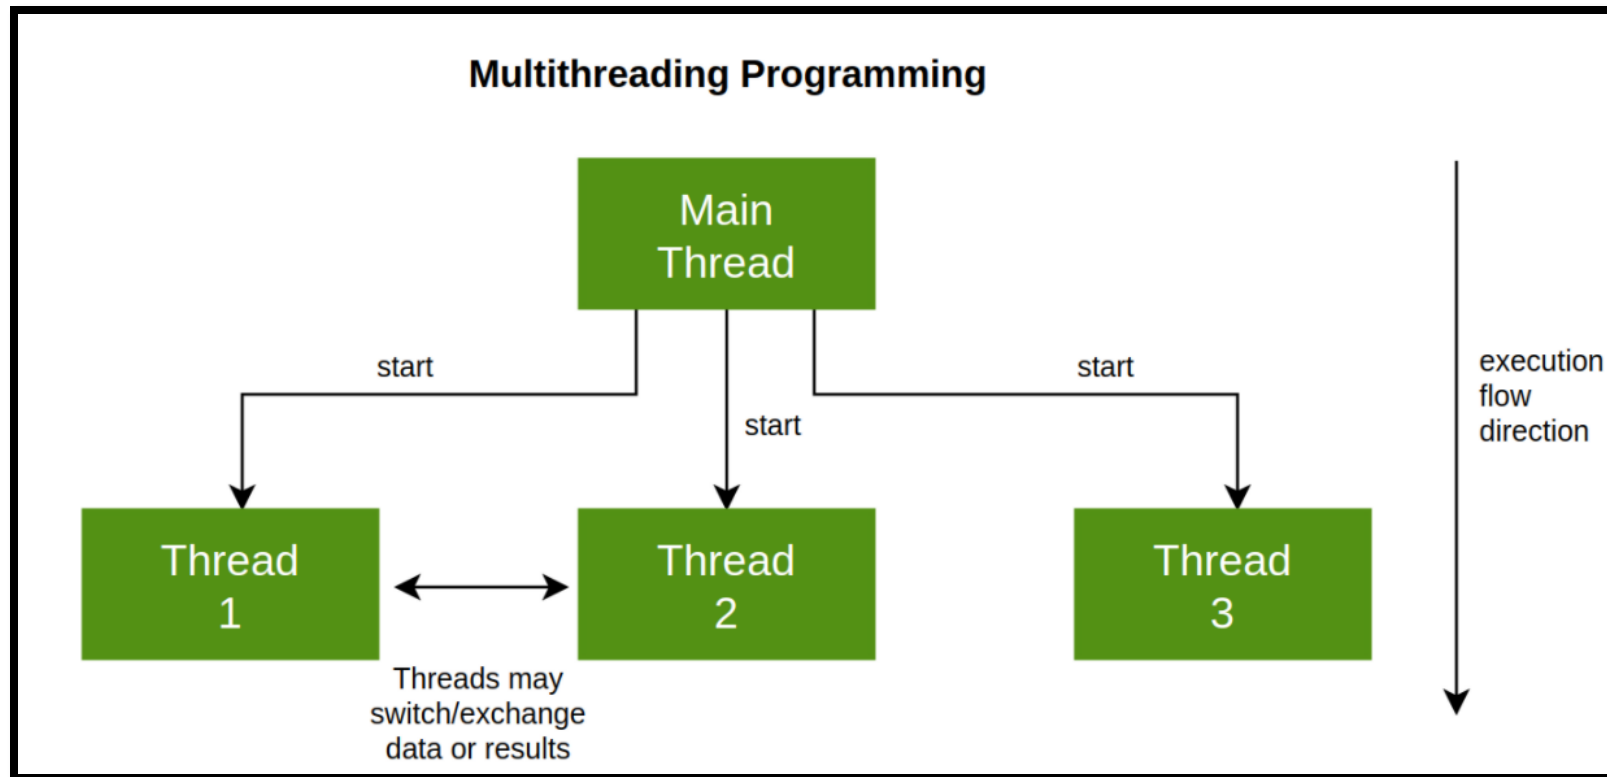
Multi-threaded process

# Why Multithreading?

Threads are popular way to improve application through parallelism. For example, in a browser, multiple tabs can be different threads.

Threads operate faster than processes due to following reasons:

- 1) Thread creation is much faster.
- 2) Context switching between threads is much faster.
- 3) Threads can be terminated easily
- 4) Communication between threads is faster.



## Can we write multithreading programs in C?


Unlike Java, multithreading is not supported by the C language standard. [POSIX Threads \(or Pthreads\)](#) is a POSIX standard for threads. Implementation of pthread is available with gcc compiler.

- To the software developer, the concept of a "procedure" that runs independently from its main program may best describe a thread.
- Pthread: POSIX Thread, a standard-based thread API for C.
- other options: openMP, std::thread
- When compiling Pthread in gcc/g++, should add option “-lpthread”.
  - Compile: `gcc test.c -lpthread` or `g++ test.cpp -lpthread`
  - Execution: `./a.out`

# Pthread creation

- `pthread_create`:


- `int pthread_create( pthread_t *thread,  
const pthread_attr_t *attr,  
void *(*start_routine) (void *),  
void *arg);`



- **thread**: pointer to an unsigned integer value that returns the thread id of the thread created.
- The **attr** parameter is used to set thread attributes. You can specify a thread attributes object like scheduling policy, detached state, etc. Set NULL by default.
- The **start\_routine** is the C routine that the thread will execute once it is created.
- **arg**: pointer to void that contains the arguments to the function defined in the earlier argument

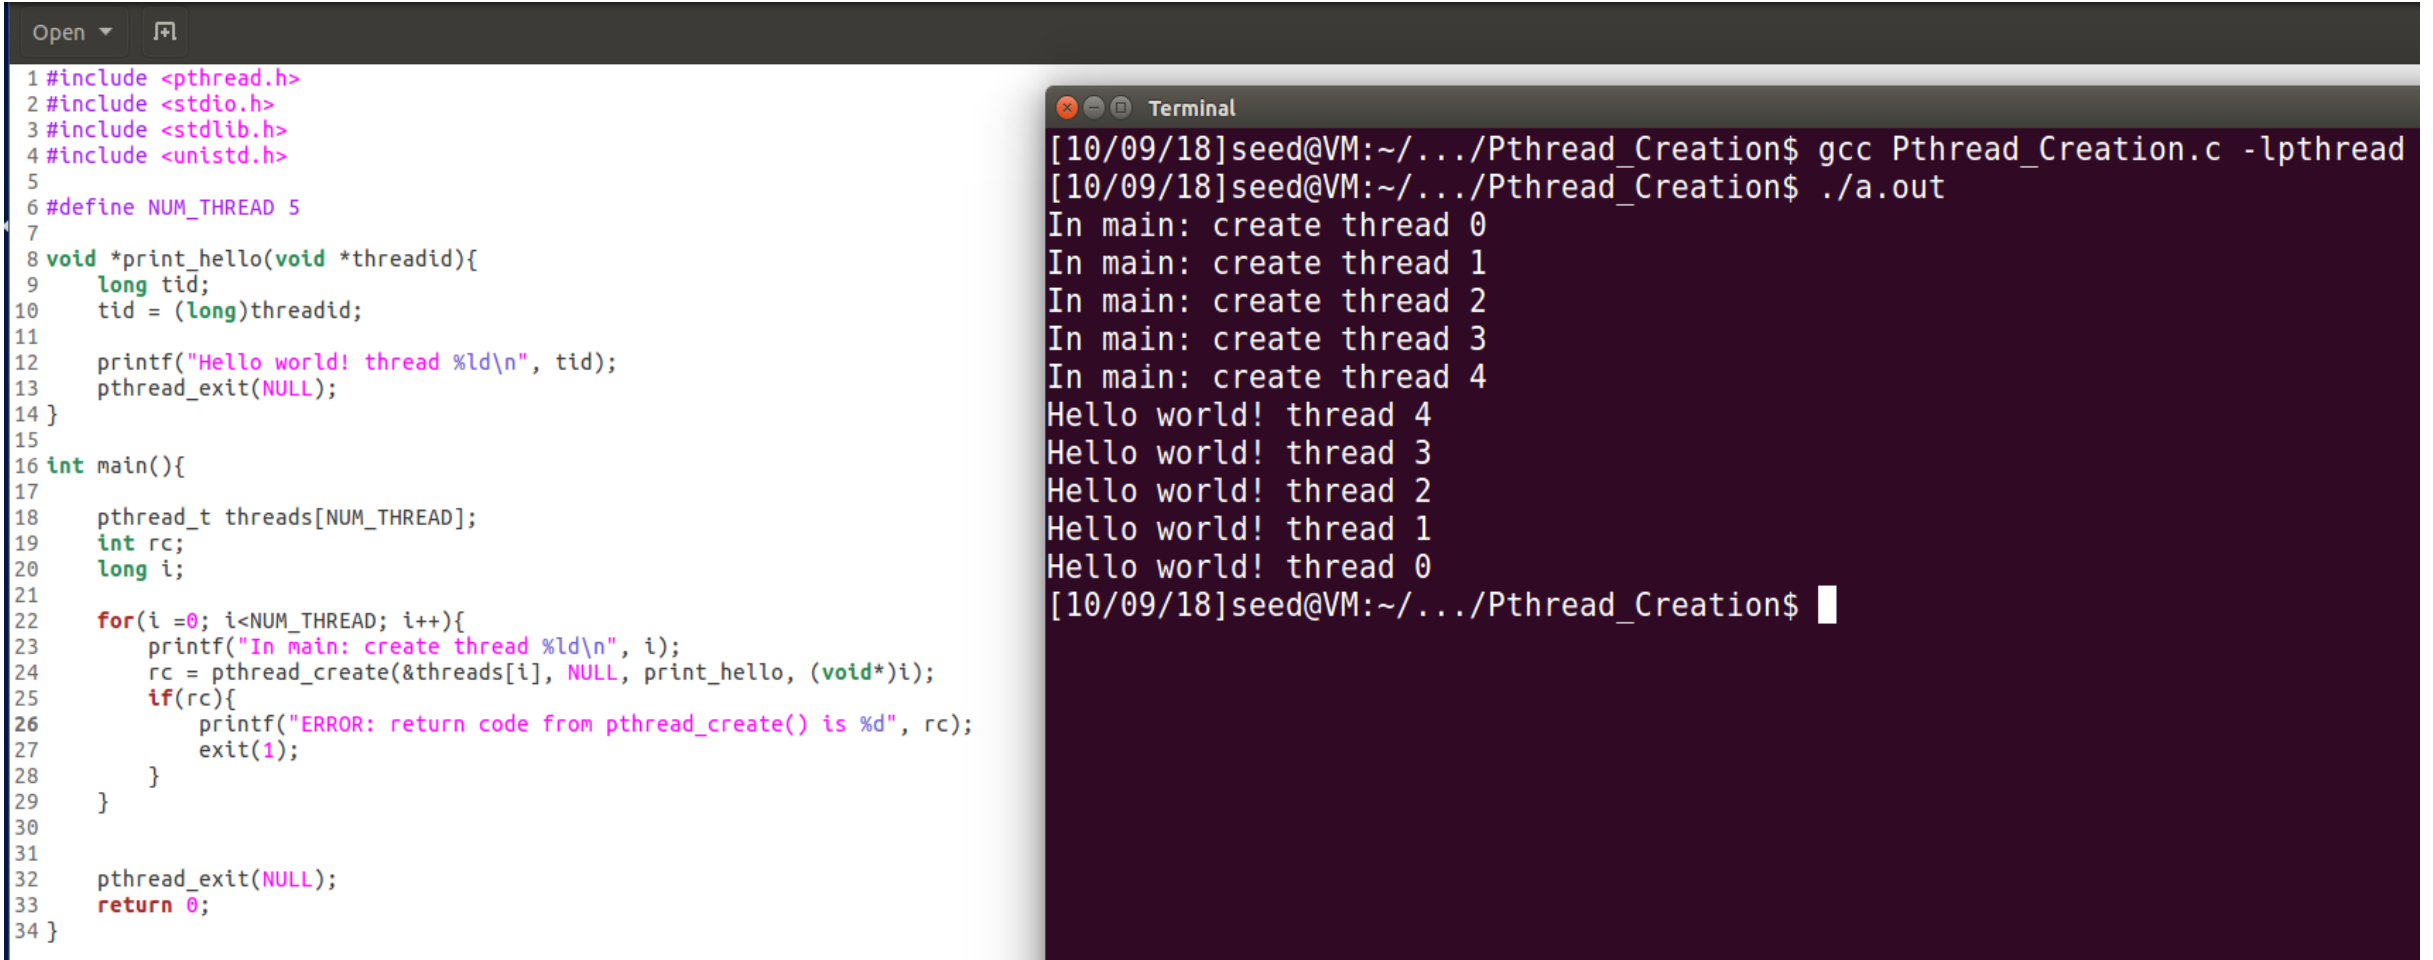
# Pthread creation

- Return value
  - On success, `pthread_create()` returns 0;
  - On error, it returns an error number, and the contents of `*thread` are undefined.
- Pthread is declared with type:
  - `pthread_t` (defined in “`sys/types.h`”)





# Pthread creation

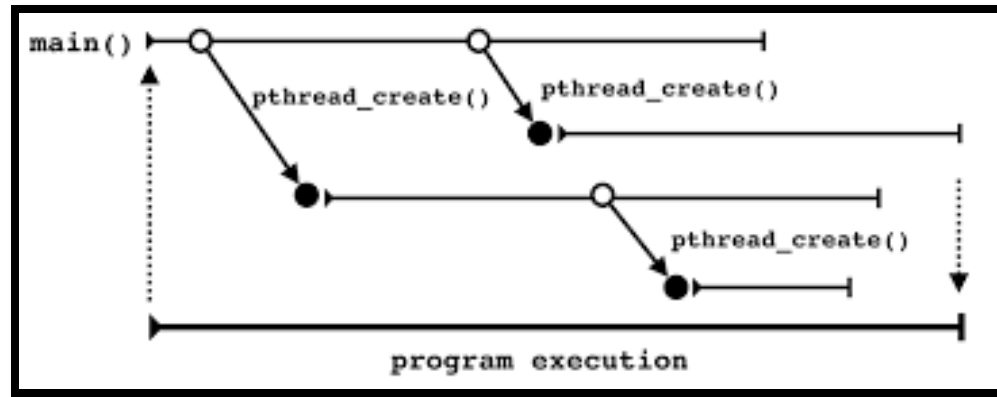


The image shows a code editor on the left and a terminal window on the right. The code editor displays the source code for a C program that creates 5 pthreads. The terminal window shows the compilation and execution of this program, with output messages from the main thread and the 5 worker threads.

```
1 #include <pthread.h>
2 #include <stdio.h>
3 #include <stdlib.h>
4 #include <unistd.h>
5
6 #define NUM_THREAD 5
7
8 void *print_hello(void *threadid){
9     long tid;
10    tid = (long)threadid;
11
12    printf("Hello world! thread %ld\n", tid);
13    pthread_exit(NULL);
14 }
15
16 int main(){
17     pthread_t threads[NUM_THREAD];
18     int rc;
19     long i;
20
21     for(i=0; i<NUM_THREAD; i++){
22         printf("In main: create thread %ld\n", i);
23         rc = pthread_create(&threads[i], NULL, print_hello, (void*)i);
24         if(rc){
25             printf("ERROR: return code from pthread_create() is %d", rc);
26             exit(1);
27         }
28     }
29
30     pthread_exit(NULL);
31     return 0;
32 }
```

```
Terminal
[10/09/18]seed@VM:~/.../Pthread_Creation$ gcc Pthread_Creation.c -lpthread
[10/09/18]seed@VM:~/.../Pthread_Creation$ ./a.out
In main: create thread 0
In main: create thread 1
In main: create thread 2
In main: create thread 3
In main: create thread 4
Hello world! thread 4
Hello world! thread 3
Hello world! thread 2
Hello world! thread 1
Hello world! thread 0
[10/09/18]seed@VM:~/.../Pthread_Creation$
```

## Pthread termination



- pthread\_exit:

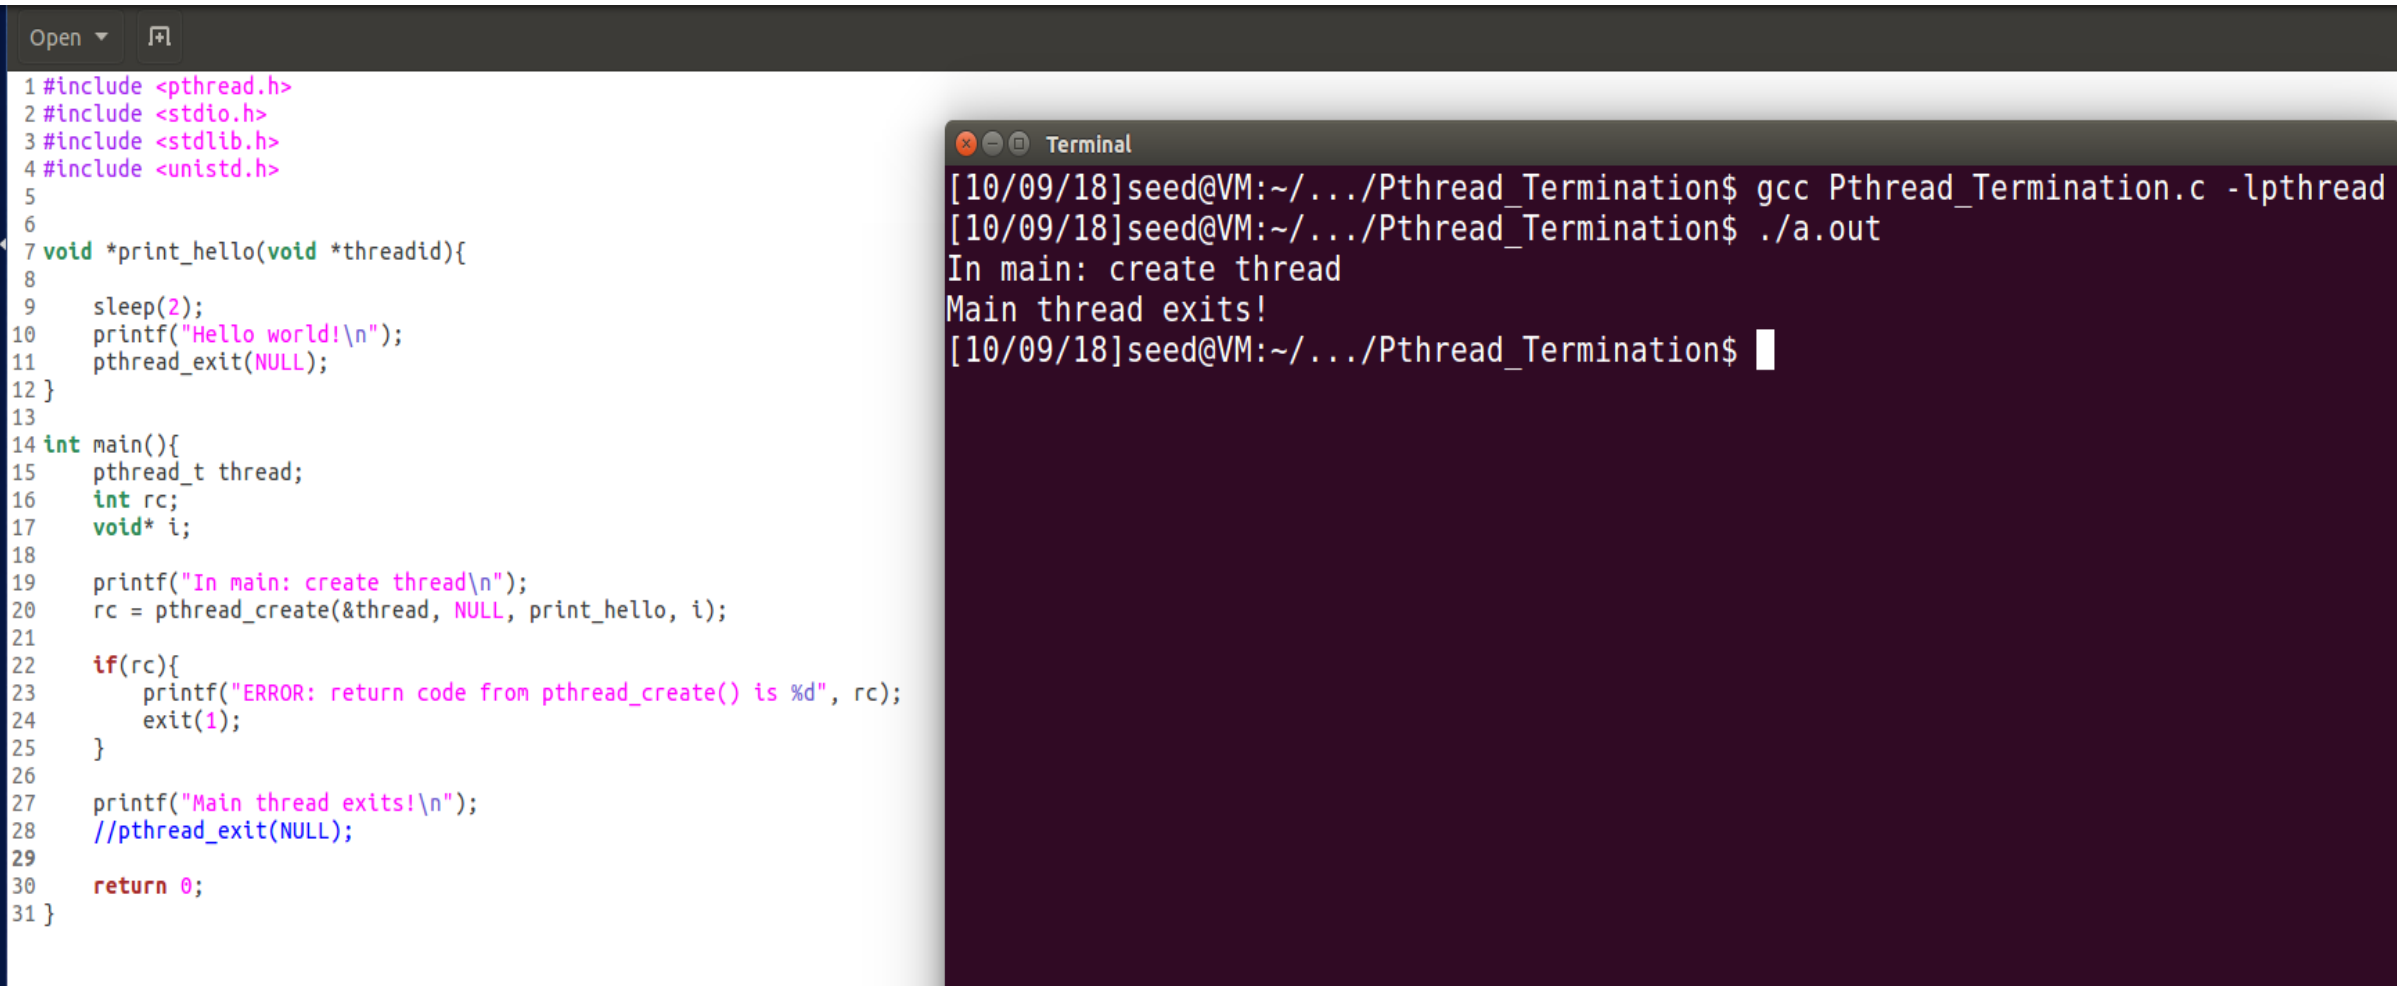
- `void pthread_exit(void *retval);`

**Parameters:** This method accepts a mandatory parameter **retval** which is the pointer to an integer that stores the return status of the thread terminated. The scope of this variable must be **global** so that any thread waiting to join this thread may read the return status.

- This routine is used to **explicitly** exit a thread. Typically, the pthread\_exit() routine is called after a thread has completed its work.
- If main() finishes before the **threads it has created, and exits with pthread\_exit()**, the other threads will **continue** to execute. Otherwise, they will **be automatically terminated** when main() finishes.
- Recommendation: **Use pthread\_exit() to exit from all threads...especially main().**

## Pthread termination – main thread without `pthread_exit()`

**Other threads would be automatically terminated** when `main()` finishes. Therefore, they might not have finished their work.



```
1 #include <pthread.h>
2 #include <stdio.h>
3 #include <stdlib.h>
4 #include <unistd.h>
5
6
7 void *print_hello(void *threadid){
8
9     sleep(2);
10    printf("Hello world!\n");
11    pthread_exit(NULL);
12 }
13
14 int main(){
15     pthread_t thread;
16     int rc;
17     void* i;
18
19     printf("In main: create thread\n");
20     rc = pthread_create(&thread, NULL, print_hello, i);
21
22     if(rc){
23         printf("ERROR: return code from pthread_create() is %d", rc);
24         exit(1);
25     }
26
27     printf("Main thread exits!\n");
28     //pthread_exit(NULL);
29
30     return 0;
31 }
```

```
Terminal
[10/09/18]seed@VM:~/.../Pthread_Termination$ gcc Pthread_Termination.c -lpthread
[10/09/18]seed@VM:~/.../Pthread_Termination$ ./a.out
In main: create thread
Main thread exits!
[10/09/18]seed@VM:~/.../Pthread_Termination$
```

Pthread termination – main with `pthread_exit()`.

The other threads will `continue` to execute.

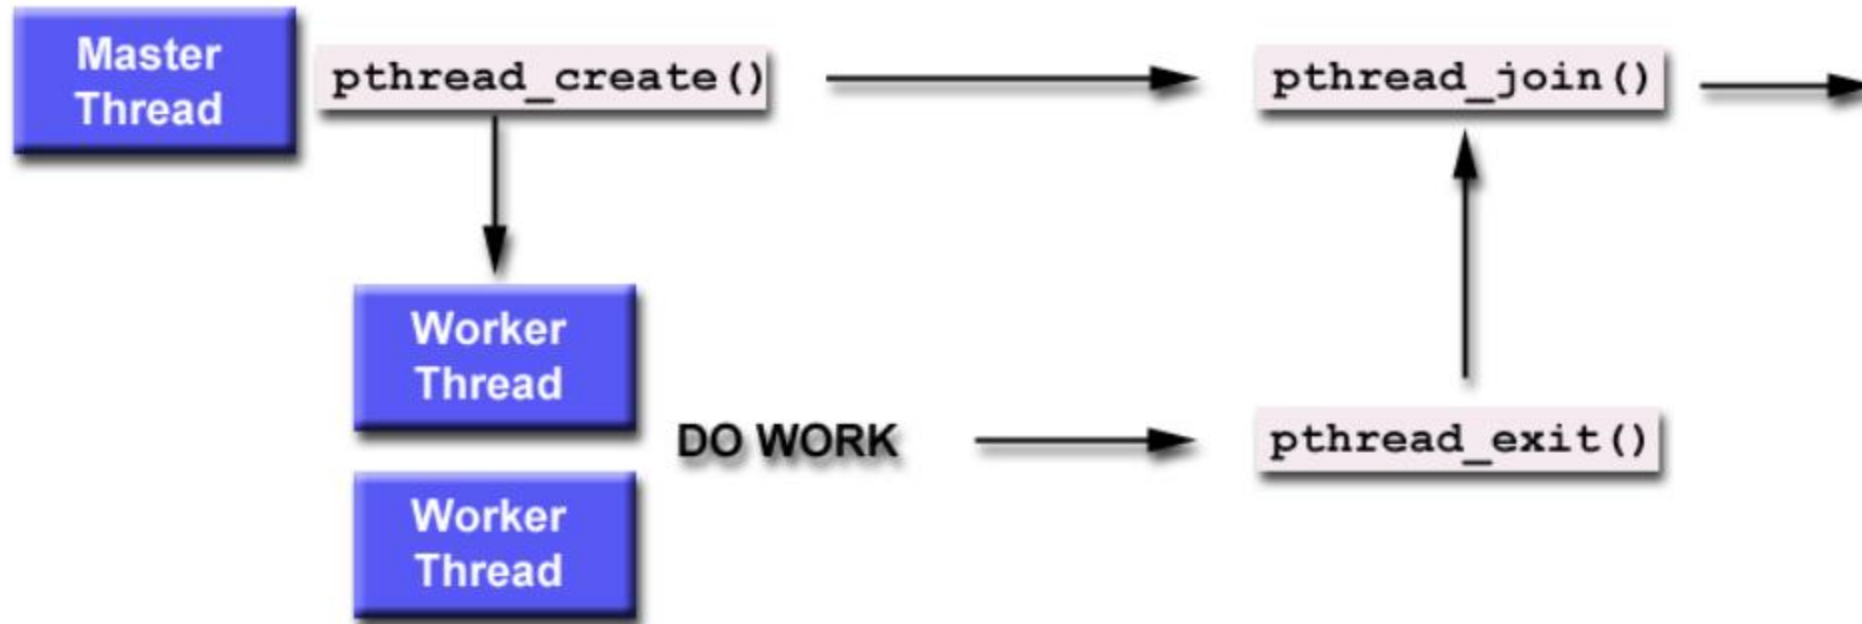
```
1 #include <pthread.h>
2 #include <stdio.h>
3 #include <stdlib.h>
4 #include <unistd.h>
5
6
7 void *print_hello(void *threadid){
8
9     sleep(2);
10    printf("Hello world!\n");
11    pthread_exit(NULL);
12 }
13
14 int main(){
15     pthread_t thread;
16     int rc;
17     void* i;
18
19     printf("In main: create thread\n");
20     rc = pthread_create(&thread, NULL, print_hello, i);
21
22     if(rc){
23         printf("ERROR: return code from pthread_create() is %d", rc);
24         exit(1);
25     }
26
27     printf("Main thread exits!\n");
28     pthread_exit(NULL);
29
30     return 0;
31 }
```

```
Terminal
[10/08/18]seed@VM:~/.../Pthread_Termination$ gcc Pthread_Termination.c -lpthread
[10/08/18]seed@VM:~/.../Pthread_Termination$ ./a.out
In main: create thread
Main thread exits!
Hello world!
[10/08/18]seed@VM:~/.../Pthread_Termination$
```

# Pthread join - synchronization

- pthread\_join:
  - `int pthread_join(pthread_t thread, void *retval);`
- "Joining" is one way to accomplish ***synchronization*** between threads.
- The pthread\_join() subroutine **blocks the calling thread** until the **specified thread terminates**.
- The programmer is able to obtain the target thread's termination return status if specified through pthread\_exit(), in the status parameter.


"Joining" is one way to accomplish synchronization between threads.



# Pthread join - synchronization

- Return value
  - On success, pthread\_join() returns **0**;
  - On error, it returns an **error number**.
- It is impossible to join a detached thread.
- When a thread is created, one of its attributes defines whether it is joinable or detached. Detached means it can never be joined. (**PTHREAD\_CREATE\_DETACHED** or **PTHREAD\_CREATE\_JOINABLE**)

# Pthread join – without calling pthread\_join().

```
Open ▾ 
1 #include<stdlib.h>
2 #include<stdio.h>
3 #include<unistd.h>
4 #include<pthread.h>
5
6
7 int sum;
8
9 void * add1(void *cnt)
10 {
11     for(int i=0; i < 5; i++)
12     {
13         sum += i;
14     }
15     pthread_exit(NULL);
16     return 0;
17 }
18 void * add2(void *cnt)
19 {
20
21     for(int i=5; i<10; i++)
22     {
23         sum += i;
24     }
25     pthread_exit(NULL);
26     return 0;
27 }
28
29 int main(void)
30 {
31     pthread_t ptid1, ptid2;
32     sum=0;
33
34     pthread_create(&ptid1, NULL, add1, &sum);
35     pthread_create(&ptid2, NULL, add2, &sum);
36
37     //pthread_join(ptid1,NULL);
38     //pthread_join(ptid2,NULL);
39
40     printf("sum %d\n", sum);
41     pthread_exit(NULL);
42
43
44     return 0;
45 }
```

```
Terminal
[10/08/18]seed@VM:~/.../Pthread_Join$ gcc Pthread_Join.c -lpthread
[10/08/18]seed@VM:~/.../Pthread_Join$ ./a.out
sum 0
[10/08/18]seed@VM:~/.../Pthread_Join$
```



## Pthread join – synchronization with calling pthread\_join().

```
Open ▾ 
1 #include<stdlib.h>
2 #include<stdio.h>
3 #include<unistd.h>
4 #include<pthread.h>
5
6
7 int sum;
8
9 void * add1(void *cnt)
10 {
11     for(int i=0; i < 5; i++)
12     {
13         sum += i;
14     }
15     pthread_exit(NULL);
16     return 0;
17 }
18 void * add2(void *cnt)
19 {
20
21     for(int i=5; i<10; i++)
22     {
23         sum += i;
24     }
25     pthread_exit(NULL);
26     return 0;
27 }
28
29 int main(void)
30 {
31     pthread_t ptid1, ptid2;
32     sum=0;
33
34     pthread_create(&ptid1, NULL, add1, &sum);
35     pthread_create(&ptid2, NULL, add2, &sum);
36
37     pthread_join(ptid1,NULL);
38     pthread_join(ptid2,NULL);
39
40     printf("sum %d\n", sum);
41     pthread_exit(NULL);
42
43
44     return 0;
45 }
```

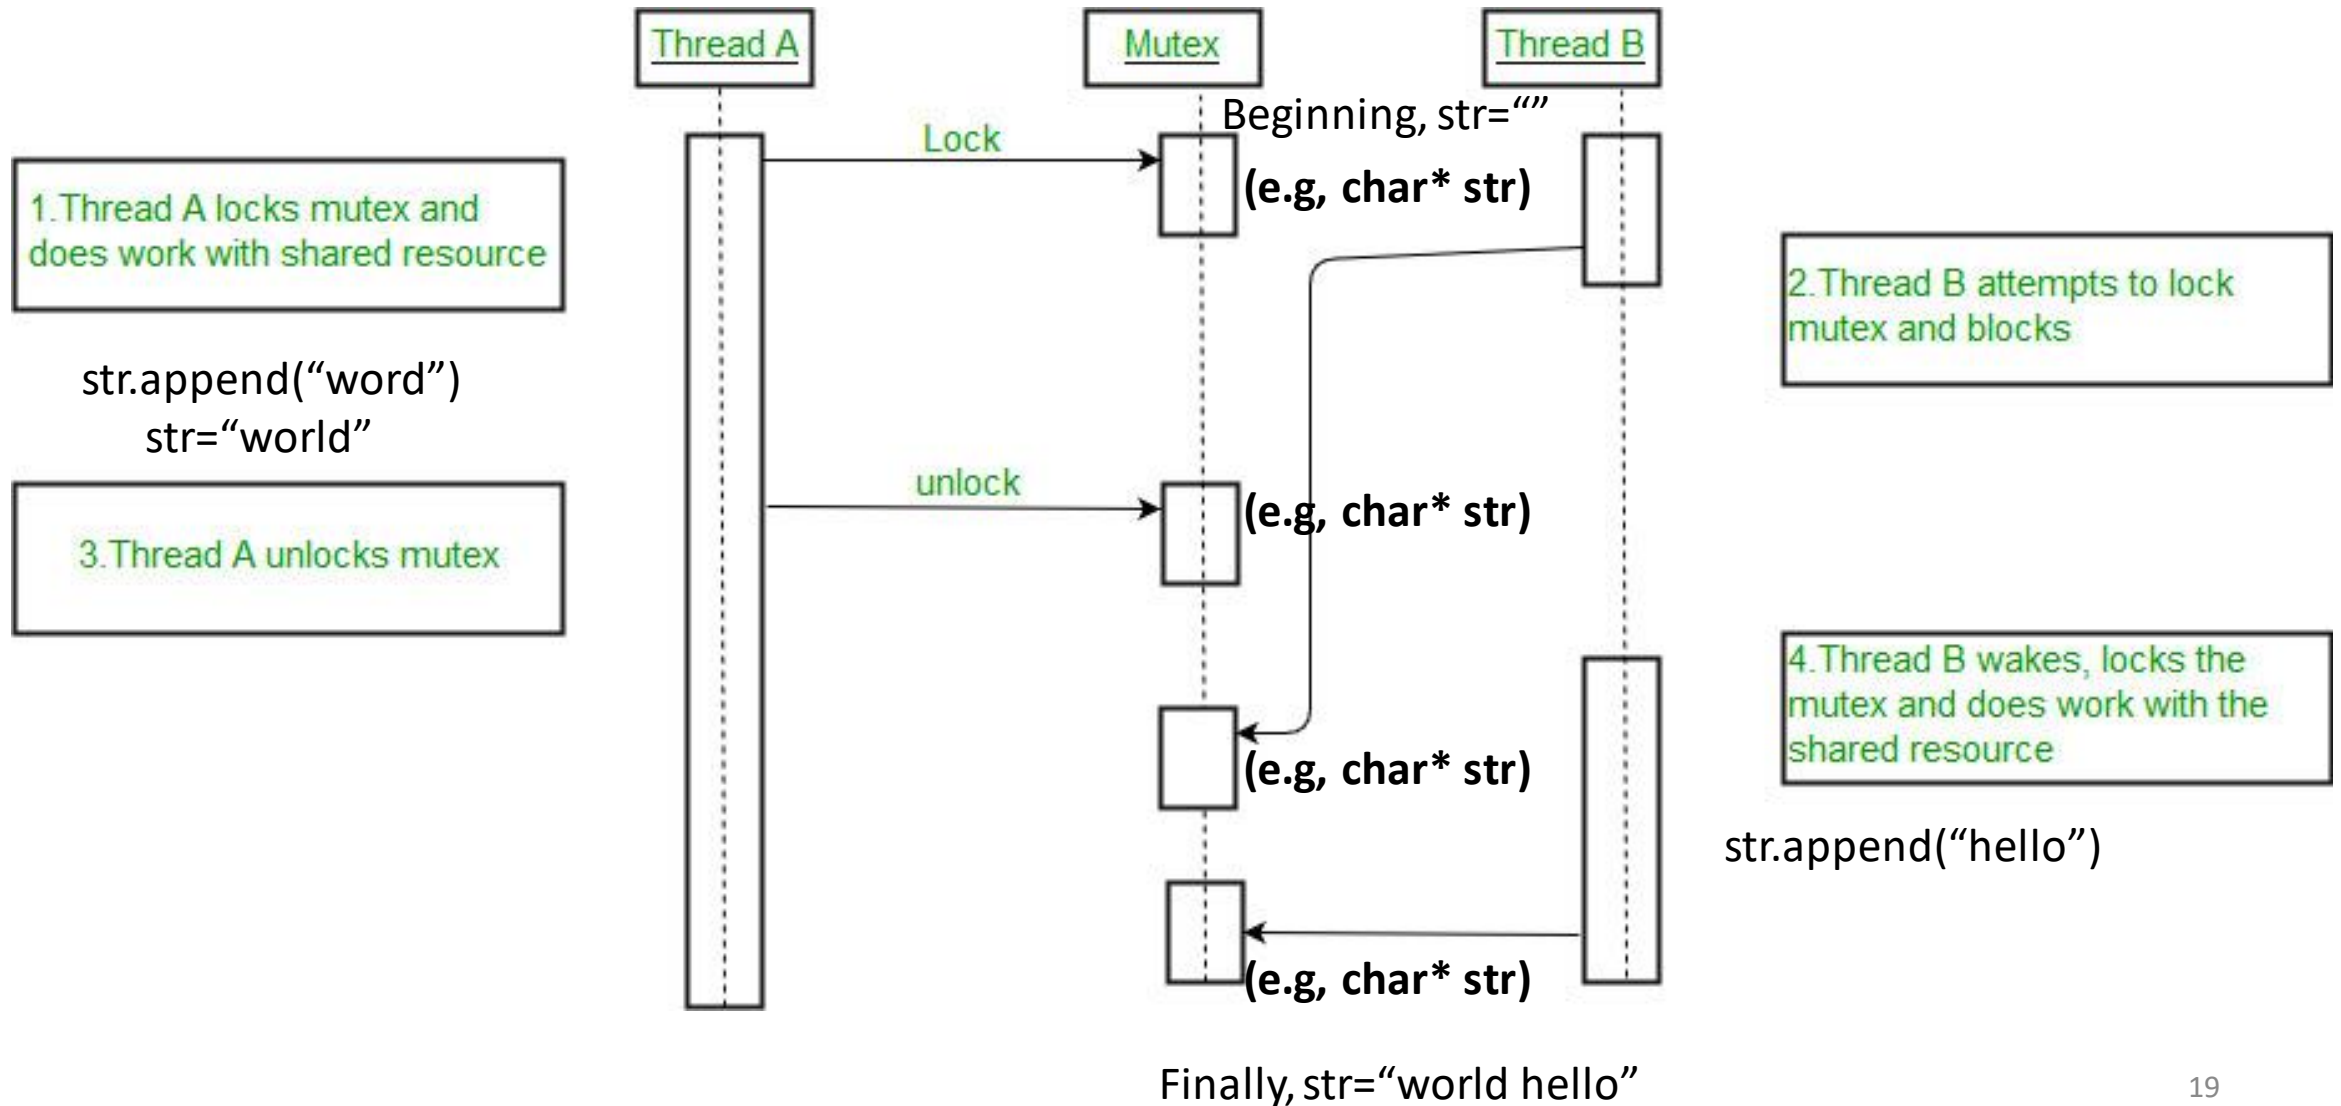
```
Terminal
[10/08/18]seed@VM:~/.../Pthread_Join$ gcc Pthread_Join.c -lpthread
[10/08/18]seed@VM:~/.../Pthread_Join$ ./a.out
sum 45
[10/08/18]seed@VM:~/.../Pthread_Join$
```

The main thread would stop and wait for ptid1 and ptid2 to finish at line 38.

# Pthread mutex – flag for privacy/security

- Mutex is an abbreviation for "**mutual exclusion**". Mutex variables are one of the primary means of implementing **thread synchronization** and for protecting shared data when multiple writes occur.
- **A mutex variable acts like a "lock" protecting access to a shared data resource.**

- A mutex variable acts like a "lock" protecting access to a shared data resource.



# Pthread mutex - flag

- Mutex should be declared with type:
  - `pthread_mutex_t` (defined in “sys/types.h”)
- Mutex should be initialized before it is used:
  - `int pthread_mutex_init( pthread_mutex_t *mutex, const pthread_mutexattr_t *attr);`
  - It initialises the `mutex` referenced by `mutex` with attributes specified by `attr`.
  - If `attr` is NULL, the default mutex attributes are used; the effect is the same as passing the address of a default mutex attributes object.
  - Upon successful initialisation, the state of the mutex becomes initialised and unlocked.
- Mutex should be free if it is no longer used:
  - `int pthread_mutex_destroy(pthread_mutex_t *);`

# Pthread mutex - flag

- Pthread mutex lock routines:

- `int pthread_mutex_lock(pthread_mutex_t *mutex);`
- `int pthread_mutex_trylock(pthread_mutex_t *mutex);`
- `int pthread_mutex_unlock(pthread_mutex_t *mutex);`

# Pthread mutex - flag

- The **pthread\_mutex\_lock()** routine is used by a thread **to acquire a lock** on the specified mutex variable. **If the mutex is already locked by another thread, this call will block the calling thread until the mutex is unlocked.**
- **pthread\_mutex\_trylock()** will attempt to lock a mutex. However, if the mutex is already locked, the routine will return immediately with a "busy" error code. This routine may be useful in preventing deadlock conditions, as in a priority-inversion situation.
- **pthread\_mutex\_unlock()** will unlock a mutex if called by the owning thread. Calling this routine is required after a thread has completed its use of protected data if other threads are to acquire the mutex for their work with the protected data. An **error** will be returned if:
  - If the mutex was already unlocked
  - If the mutex is owned by another thread

Line 20 consume more time (e.g. 3seconds)

### Execution Order: A B C D E F

Thread 1 (execute first)

Thread 2

A Line 17: counter=1

C Line 17: counter=2

B Line 20: much time

D Line 20: much time

E Line 22: J"Job 2"

F Line 22: J"Job 2"

```
Run: pthread_without_mutex x
/home/lemaker/open-source/CLionProjects/f:

Job 1 has started
Job 2 has started
Job 2 has finished
Job 2 has finished
Process finished with exit code 0
```

```

13 void* trythis(void* arg)
14 {
15     pthread_mutex_lock(&lock);
16     unsigned long i = 0;
17     counter += 1;
18     printf( format: "\n Job %d has started\n", counter);
19
20     for (i = 0; i < (0xFFFFFFFF); i++)
21         ;
22     printf( format: "\n Job %d has finished\n", counter);
23
24     pthread_mutex_unlock(&lock);
25
26     return NULL;
27 }

```

```

33 int main(void)
34 {
35     int i = 0;
36     int error;
37     counter=0;
38     if (pthread_mutex_init(&lock, mutexattr: NULL) != 0) {
39         printf( format: "\n mutex init has failed\n");
40         return 1;
41     }
42     while (i < 2) {
43         error = pthread_create(&(tid[i]),
44                               attr: NULL,
45                               &trythis, arg: NULL);
46         if (error != 0)
47             printf( format: "\nThread can't be created :[%s]",
48                     strerror(error));
49         i++;
50     }
51
52     pthread_join( th: tid[0], thread_return: NULL);
53     pthread_join( th: tid[1], thread_return: NULL);
54     pthread_mutex_destroy(&lock);

```

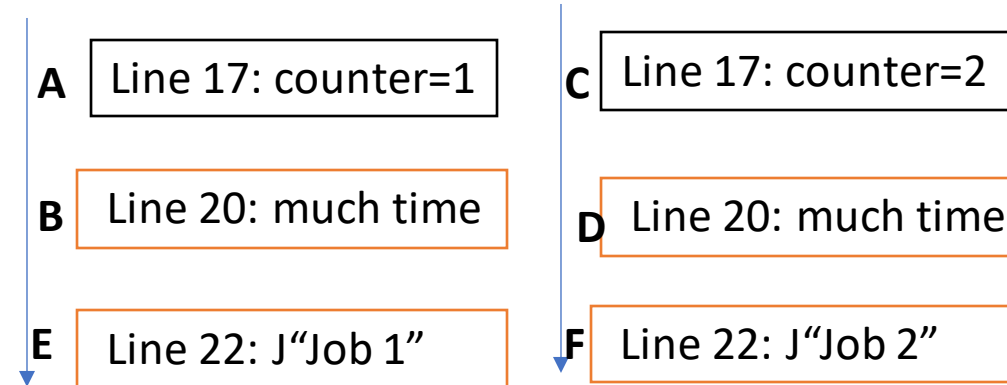
Line 20 consume more time (e.g. 3seconds)

**Execution Order: A B C D E F (impossible)**

**Execution Order: A B E C D F or C D F A B E**

Thread 1 (execute first)

Thread 2



```

Job 1 has started

Job 1 has finished

Job 2 has started

Job 2 has finished

Process finished with exit code 0

```



# Pthread condition - signals

- **Condition variables** provide yet another way for threads to synchronize.
- While **mutexes** implement synchronization by controlling thread access to data, condition variables allow threads **to synchronize based upon the actual value of data**.
- A condition variable is always **used in conjunction with a mutex lock**.

# Pthread condition - signals

- Condition variables must be declared with type: `pthread_cond_t`
  - `pthread_cond_t` (defined in “sys/types.h”)
- Condition variables must be initialized before it is used:
  - `int pthread_cond_init(pthread_cond_t *, const pthread_condattr_t *);`
- Condition variables should be freed if it is no longer used:
  - `int pthread_cond_destroy(pthread_cond_t *);`

# Pthread condition - signals

- Pthread condition routines:

- `int pthread_cond_wait(pthread_cond_t *, pthread_mutex_t *);`
- `int pthread_cond_signal(pthread_cond_t *);`
- `int pthread_cond_broadcast(pthread_cond_t *);`

# Pthread condition - signals

- **pthread\_cond\_wait()** blocks the calling thread until **the specified condition** is signalled. This routine should be called **while mutex is locked**, and it will **automatically release the mutex while it waits**.
- **The pthread\_cond\_signal()** routine is used **to signal (or wake up) another thread which is waiting** on the condition variable. It should be called after mutex is locked, and must unlock mutex in order for pthread\_cond\_wait() routine to complete.
- **The pthread\_cond\_broadcast()** routine should be used instead of pthread\_cond\_signal() if more than one thread is in a blocking wait state.

# Pthread condition - signals

```
Open ▾ [F]
1 #include <pthread.h>
2 #include <stdio.h>
3 #include <unistd.h>
4
5 #define NUM_THREADS 3
6 #define TCOUNT 10
7 #define COUNT_LIMIT 10
8
9 int count = 0;
10 int thread_ids[3] = {0,1,2};
11 pthread_mutex_t count_mutex;
12 pthread_cond_t count_threshold_cv;
13
14 void *inc_count(void *idp)
15 {
16     int i = 0;
17     int taskid = 0;
18     int *my_id = (int*)idp;
19
20     for (i=0; i<TCOUNT; i++) {
21         pthread_mutex_lock(&count_mutex);
22         taskid = count;
23         count++;
24
25         if (count == COUNT_LIMIT){
26             pthread_cond_signal(&count_threshold_cv);
27         }
28
29         printf("inc_count(): thread %d, count = %d, unlocking mutex\n", *my_id, count);
30         pthread_mutex_unlock(&count_mutex);
31         sleep(1);
32     }
33
34     printf("inc_count(): thread %d, Threshold reached.\n", *my_id);
35
36     pthread_exit(NULL);
37 }
--
```

```
39 void *watch_count(void *idp)
40 {
41     int *my_id = (int*)idp;
42     printf("Starting watch_count(): thread %d\n", *my_id);
43
44     pthread_mutex_lock(&count_mutex);
45
46     while(count<COUNT_LIMIT) {
47         pthread_cond_wait(&count_threshold_cv, &count_mutex);
48         printf("watch_count(): thread %d Condition signal received.\n", *my_id);
49     }
50
51     count += 100;
52     pthread_mutex_unlock(&count_mutex);
53     pthread_exit(NULL);
54 }
55
56 int main (int argc, char *argv[])
57 {
58     int i, rc;
59     pthread_t threads[3];
60     pthread_attr_t attr;
61
62     /* Initialize mutex and condition variable objects */
63     pthread_mutex_init(&count_mutex, NULL);
64     pthread_cond_init (&count_threshold_cv, NULL);
65
66     /* For portability, explicitly create threads in a joinable state */
67     pthread_attr_init(&attr);
68     pthread_attr_setdetachstate(&attr, PTHREAD_CREATE_JOINABLE);
69     pthread_create(&threads[0], &attr, inc_count, (void *)&thread_ids[0]);
70     pthread_create(&threads[1], &attr, inc_count, (void *)&thread_ids[1]);
71     pthread_create(&threads[2], &attr, watch_count, (void *)&thread_ids[2]);
72
73     /* Wait for all threads to complete */
74     for (i=0; i<NUM_THREADS; i++) {
75         pthread_join(threads[i], NULL);
76     }
77     printf ("Main(): Waited on %d threads. Done.\n", NUM_THREADS);
78
79     /* Clean up and exit */
80     pthread_attr_destroy(&attr);
81     pthread_mutex_destroy(&count_mutex);
82     pthread_cond_destroy(&count_threshold_cv);
83     pthread_exit(NULL);
84
85     return 0;
86 }
```

Release the mutex at line 47



## Pthread condition - signals

```
Terminal
[10/09/18]seed@VM:~/.../Pthread_Cond$ gcc Pthread_Cond.c -lpthread
[10/09/18]seed@VM:~/.../Pthread_Cond$ ./a.out
Starting watch_count(): thread 2
inc_count(): thread 1, count = 1, unlocking mutex
inc_count(): thread 0, count = 2, unlocking mutex
inc_count(): thread 1, count = 3, unlocking mutex
inc_count(): thread 0, count = 4, unlocking mutex
inc_count(): thread 1, count = 5, unlocking mutex
inc_count(): thread 0, count = 6, unlocking mutex
inc_count(): thread 1, count = 7, unlocking mutex
inc_count(): thread 0, count = 8, unlocking mutex
inc_count(): thread 1, count = 9, unlocking mutex
inc_count(): thread 0, count = 10, unlocking mutex
watch_count(): thread 2 Condition signal received.
inc_count(): thread 1, count = 111, unlocking mutex
inc_count(): thread 0, count = 112, unlocking mutex
inc_count(): thread 1, count = 113, unlocking mutex
inc_count(): thread 0, count = 114, unlocking mutex
inc_count(): thread 1, count = 115, unlocking mutex
inc_count(): thread 0, count = 116, unlocking mutex
inc_count(): thread 1, count = 117, unlocking mutex
inc_count(): thread 0, count = 118, unlocking mutex
inc_count(): thread 1, count = 119, unlocking mutex
inc_count(): thread 0, count = 120, unlocking mutex
inc_count(): thread 1, Threshold reached.
inc_count(): thread 0, Threshold reached.
Main(): Waited on 3 threads. Done.
[10/09/18]seed@VM:~/.../Pthread_Cond$
```

# References

- <https://www.baeldung.com/cs/async-vs-multi-threading>
- <https://www.javatpoint.com/process-vs-thread>
- <https://www.geeksforgeeks.org/multithreading-c-2/>
- <https://www.geeksforgeeks.org/condition-wait-signal-multi-threading/>
- [http://www.cs.unibo.it/~ghini/didattica/sistop/pthreads\\_tutorial/POSIX\\_Threads\\_Programming.htm](http://www.cs.unibo.it/~ghini/didattica/sistop/pthreads_tutorial/POSIX_Threads_Programming.htm)
- <https://www.geeksforgeeks.org/mutex-lock-for-linux-thread-synchronization/>

Thank you